

Untitled

October 11, 2018

```
In [2]: import numpy
import pandas
```

```
In [3]: train_dataframe = pandas.read_csv("./train.csv")
test_dataframe = pandas.read_csv("./test.csv")
```

```
In [4]: train_dataframe.head()
```

```
Out[4]:
```

	age	gender	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [5]: test_dataframe.head()
```

```
Out[5]:
```

	age	gender	bmi	children	smoker	region	charges
0	30	male	22.99	2	yes	northwest	17361.76610
1	24	male	32.70	0	yes	southwest	34472.84100
2	24	male	25.80	0	no	southwest	1972.95000
3	48	male	29.60	0	no	southwest	21232.18226
4	47	male	19.19	1	no	northeast	8627.54110

```
In [6]: train_dataframe.dtypes
```

```
Out[6]: age          int64
gender         object
bmi           float64
children       int64
smoker         object
region         object
charges       float64
dtype: object
```

```
In [7]: train_dataframe["gender"].value_counts()
```

```
Out[7]: male          505
female         495
Name: gender, dtype: int64
```

```
In [8]: cleanup_nums = {"gender":{"male": 1, "female": 0}}
```

```
In [9]: train_dataframe.replace(cleanup_nums, inplace=True)
test_dataframe.replace(cleanup_nums, inplace=True)
train_dataframe.head()
```

```
Out[9]:
```

	age	gender	bmi	children	smoker	region	charges
0	19	0	27.900	0	yes	southwest	16884.92400
1	18	1	33.770	1	no	southeast	1725.55230
2	28	1	33.000	3	no	southeast	4449.46200
3	33	1	22.705	0	no	northwest	21984.47061
4	32	1	28.880	0	no	northwest	3866.85520

```
In [10]: test_dataframe.head()
```

```
Out[10]:
```

	age	gender	bmi	children	smoker	region	charges
0	30	1	22.99	2	yes	northwest	17361.76610
1	24	1	32.70	0	yes	southwest	34472.84100
2	24	1	25.80	0	no	southwest	1972.95000
3	48	1	29.60	0	no	southwest	21232.18226
4	47	1	19.19	1	no	northeast	8627.54110

```
In [11]: train_dataframe["gender"].value_counts()
```

```
Out[11]: 1    505
         0    495
         Name: gender, dtype: int64
```

```
In [12]: cleanup_nums = {"smoker":{"yes": 1, "no": 0}}
```

```
In [13]: train_dataframe.replace(cleanup_nums, inplace=True)
test_dataframe.replace(cleanup_nums, inplace=True)
train_dataframe.head()
```

```
Out[13]:
```

	age	gender	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	southwest	16884.92400
1	18	1	33.770	1	0	southeast	1725.55230
2	28	1	33.000	3	0	southeast	4449.46200
3	33	1	22.705	0	0	northwest	21984.47061
4	32	1	28.880	0	0	northwest	3866.85520

```
In [14]: train_dataframe["region"].value_counts()
```

```
Out[14]: southeast    278
         northeast    247
         southwest    244
         northwest    231
         Name: region, dtype: int64
```

```
In [15]: train_dataframe = pandas.get_dummies(train_dataframe, columns=['region'])
test_dataframe = pandas.get_dummies(test_dataframe, columns=['region'])
```

```
In [16]: test_dataframe.head()
```

```
Out[16]:
```

	age	gender	bmi	children	smoker	charges	region_northeast	\
0	30	1	22.99	2	1	17361.76610	0	
1	24	1	32.70	0	1	34472.84100	0	
2	24	1	25.80	0	0	1972.95000	0	
3	48	1	29.60	0	0	21232.18226	0	
4	47	1	19.19	1	0	8627.54110	1	

	region_northwest	region_southeast	region_southwest
0	1	0	0
1	0	0	1
2	0	0	1
3	0	0	1
4	0	0	0

0.1 regression without regularization

```
In [17]: train_dataframe.describe()
```

```
Out[17]:
```

	age	gender	bmi	children	smoker	\
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	
mean	39.615000	0.505000	30.86338	1.080000	0.196000	
std	14.153908	0.500225	6.04744	1.198765	0.397167	
min	18.000000	0.000000	15.96000	0.000000	0.000000	
25%	27.000000	0.000000	26.60000	0.000000	0.000000	
50%	40.000000	1.000000	30.59000	1.000000	0.000000	
75%	52.000000	1.000000	35.11250	2.000000	0.000000	
max	64.000000	1.000000	50.38000	5.000000	1.000000	

	charges	region_northeast	region_northwest	region_southeast	\
count	1000.000000	1000.000000	1000.000000	1000.000000	
mean	13075.755883	0.247000	0.231000	0.278000	
std	11985.924552	0.431483	0.421683	0.448238	
min	1121.873900	0.000000	0.000000	0.000000	
25%	4719.683425	0.000000	0.000000	0.000000	
50%	9283.021300	0.000000	0.000000	0.000000	
75%	15882.795437	0.000000	0.000000	1.000000	
max	63770.428010	1.000000	1.000000	1.000000	

	region_southwest
count	1000.000000
mean	0.244000
std	0.429708
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
In [18]: train_dataframe['age'] = train_dataframe['age'].apply(lambda x: x ** 2)
test_dataframe['age'] = test_dataframe['age'].apply(lambda x: x ** 2)
train_dataframe.head()
```

```
Out [18]:
```

	age	gender	bmi	children	smoker	charges	region_northeast \
0	361	0	27.900	0	1	16884.92400	0
1	324	1	33.770	1	0	1725.55230	0
2	784	1	33.000	3	0	4449.46200	0
3	1089	1	22.705	0	0	21984.47061	0
4	1024	1	28.880	0	0	3866.85520	0

	region_northwest	region_southeast	region_southwest
0	0	0	1
1	0	1	0
2	0	1	0
3	1	0	0
4	1	0	0

```
In [19]: X_train = train_dataframe.loc[:, train_dataframe.columns != 'charges'].values
Y_train = train_dataframe[['charges']].values
X_test = test_dataframe.loc[:, test_dataframe.columns != 'charges'].values
Y_test = test_dataframe[['charges']].values
```

```
In [20]: def cal_w(X,Y):
W = numpy.matmul(numpy.matmul(numpy.linalg.inv(numpy.matmul(numpy.transpose(X), X))
return W
```

```
In [21]: W = cal_w(X_train,Y_train)
```

```
In [22]: W
```

```
Out [22]: array([[ 3.30752409e+00],
[-2.83185963e+02],
[ 3.37031996e+02],
[ 5.53892111e+02],
[ 2.38823566e+04],
[-7.52569883e+03],
[-7.99692026e+03],
[-8.83897937e+03],
[-8.81785138e+03]])
```

0.2 part B

```
In [23]: def cal_MSE(Y, Yhat):
return numpy.asscalar(numpy.matmul(numpy.transpose(Y - Yhat), (Y - Yhat)))/Y.shape[0]
```

```
In [24]: sizes = []
train_errors = []
test_errors = []
```

```


for size in range(20, 1000, 10):
    W = cal_w(X_train[:size], Y_train[:size])
    Yhat_train = numpy.matmul(X_train[:size], W)
    Yhat_test = numpy.matmul(X_test, W)
    sizes.append(size)
    train_errors.append(cal_MSE(Y_train[:size], Yhat_train))
    test_errors.append(cal_MSE(Y_test, Yhat_test))
    #print(cal_MSE(Y_train[:size], Yhat_train))
    #print(cal_MSE(Y_test, Yhat_test))

```

```

In [28]: import matplotlib.pyplot as plt
plt.plot(sizes, train_errors, label='train')
plt.legend()
plt.plot(sizes, test_errors, label='test')
plt.legend()
plt.ylabel('MSE')
plt.show()

```



0.3 Batch Gradient Decent

```

In [29]: d = X_train.shape[1]
W = numpy.zeros((d,1))
step_size = 0.0000000004

In [30]: for step_num in range(1, 200000):
temp = numpy.zeros((d,1))

```

```

        for i in range (0, X_train.shape[0]):
            temp = temp + ((Y_train[i]-numpy.matmul(numpy.transpose(W), X_train[i]))*X_train[i])
            W_next = W + step_size * temp
            W = W_next
        Yhat = numpy.matmul(X_train, W)
        print("final MSE after "+str(step_num)+" step: "+str(cal_MSE(Y_train, Yhat)))

final MSE after 199999 step: 122715991.56084847

```

In [31]: W

```

Out[31]: array([[ 3.18683152],
 [ 26.21011364],
 [238.96797125],
 [ 59.74107693],
 [298.53624625],
 [ 18.30637523],
 [ -9.08822966],
 [ 13.94578975],
 [-21.89711975]])

```

```

In [32]: Yhat_test = numpy.matmul(X_test, W)
        print("final MSE of test data after "+str(step_num)+" step: "+str(cal_MSE(Y_test, Yhat_test)))

final MSE of test data after 199999 step: 151051461.36231995

```

0.4 Stochastic Gradient Decent

```

In [33]: W = numpy.zeros((d,1))
        step_size = 0.0000000004

In [34]: for step_num in range(1, 20000000):
            temp = numpy.zeros((d,1))
            i = step_num%1000
            temp = temp + ((Y_train[i]-numpy.matmul(numpy.transpose(W), X_train[i]))*X_train[i])
            W_next = W + step_size * temp
            W = W_next
        Yhat = numpy.matmul(X_train, W)
        print("final MSE after "+str(step_num)+" step: "+str(cal_MSE(Y_train, Yhat)))

final MSE after 19999999 step: 124977095.7821533

```

In [35]: W

```

Out[35]: array([[ 3.38442291e+00],
 [ 5.82913927e+00],

```

```

[ 2.19888003e+02],
[ 1.27992666e+01],
[ 3.12040760e+01],
[ 3.03571344e+00],
[ 1.70050571e-01],
[ 3.60480991e+00],
[-8.79139469e-01]])

```

```

In [36]: Yhat_test = numpy.matmul(X_test, W)
         print("final MSE of test data after "+str(step_num)+" step: "+str(cal_MSE(Y_test, Yhat_

```

```

final MSE of test data after 19999999 step: 156323270.19909942

```

0.5 Regression With Regularization

```

In [37]: lambdaI = 4*numpy.identity(X_train.shape[1])
         lambdaI.shape

```

```

Out[37]: (9, 9)

```

```

In [38]: def cal_w(X,Y, lam):
         lambdaI = lam*numpy.identity(X.shape[1])
         W = numpy.matmul(numpy.matmul(numpy.linalg.inv(numpy.matmul(numpy.transpose(X), X)+
         return W

```

```

In [39]: X_train[100:200].shape

```

```

Out[39]: (100, 9)

```

```

In [40]: validation_err = []
         for p in range(-4,5,1):
             lam = 10**p
             temp = 0
             for i in range(0, 1000, 200):
                 X_train_train = numpy.delete(X_train, range(i,i+200), 0)
                 X_train_valid = X_train[i:i+200]
                 Y_train_train = numpy.delete(Y_train, range(i,i+200), 0)
                 Y_train_valid = Y_train[i:i+200]
                 W = cal_w(X_train_train,Y_train_train, lam)
                 Yhat_train_valid = numpy.matmul(X_train_valid, W)
                 temp = temp + cal_MSE(Y_train_valid, Yhat_train_valid)
             validation_err.append(temp/5)
         W = cal_w(X_train,Y_train, lam)
         Yhat_train = numpy.matmul(X_train, W)
         print(cal_MSE(Y_train, Yhat_train))

```

```

34673667.87071548

```


```

34673667.911172494

```

```
34673671.948340915
34674067.28818611
34706568.7767418
35711282.26299803
49759909.07841898
102087142.97740214
122195557.87957357
```

```
In [41]: import matplotlib.pyplot as plt
plt.plot(range(-4,5,1), validation_err)
plt.ylabel('validation errs')
plt.show()
```



```
In [42]: lam = 1e-4
W = cal_w(X_train,Y_train, lam)
Yhat_train = numpy.matmul(X_train, W)
print(cal_MSE(Y_train, Yhat_train))
Yhat_test = numpy.matmul(X_test, W)
print(cal_MSE(Y_test, Yhat_test))
```

```
34673667.87071548
43731022.13001466
```